

CREDIT UNDERWRITING & RISK PREDICTION REPORT

Vanguard Elite Risk Intelligence Pipeline

Student Name:	Anushay Ayat
Student ID:	773317
Department:	Artificial Intelligence & Data Science
Course / Portal:	SMIT Student Portal - Saylani Mass IT Training
Project Title:	Automated Financial Risk Mitigation & Predictive Underwriting

1. Executive Summary

This technical report outlines the development and performance evaluation of an end-to-end automated credit underwriting pipeline. The objective is to leverage modern machine learning architectures to predict loan approval suitability based on historical socio-demographic vectors and capital exposure metrics. By transitioning from subjective human evaluation to data-driven algorithmic scoring, financial institutions can eliminate operational biases, lower default frequencies, and process credit requests within sub-second latencies.

2. Data Profiling & Preprocessing Pipeline

The initial data matrix consisted of a raw 13-feature customer dataset. To ensure empirical validity and protect subsequent models from structural instability, a meticulous four-stage data integrity workflow was enforced:

- **Structural Integrity Audit:** Profiling using structural maps identified missing elements across critical indicators including Gender, Married, Dependents, Self_Employed, LoanAmount, Loan_Amount_Term, and Credit_History.
- **Statistical Imputation Matrix:** To maintain standard data structures without dropping valuable rows, categorical gaps were imputed using the column **Mode**, whereas missing continuous values were resolved using the column **Median**.
- **Dimensionality Trimming:** Administrative strings such as Loan_ID were dropped from the tensor flow since they offer zero predictive variance.
- **Categorical Tensor Encoding:** Multi-class text parameters (e.g., Property_Area, Education, Gender, Married, Self_Employed, and the target Loan_Status) were successfully transformed into mathematically consumable integers using Scikit-Learn's LabelEncoder protocol. Clean string bounds on string elements like Dependents (removing text signs like "+") ensured seamless data alignment.

3. Exploratory Data Analysis (EDA) Insights

Statistical distributions and multi-variate correlations revealed pivotal system drivers:

- **Loan Status Analysis:** The target variable displayed a natural class imbalance, with a higher proportion of approved loans compared to rejections, which was accommodated during model training using stratified partitioning.
- **Income vs. Loan Approval Analysis:** Continuous distributions analyzed via kernel density estimation (KDE) and comparative box plots showed heavily right-skewed ranges for both `ApplicantIncome` and `CoapplicantIncome`, signaling that raw income alone is an insufficient predictor without adjusting for requested loan magnitudes.
- **Pearson Correlation Heatmap Matrix:** Multi-collinearity analysis verified that `Credit_History` holds an overwhelming positive correlation with loan validation outcomes, serving as a primary baseline filter.

4. Advanced Feature Engineering

To better isolate capital leverage dynamics, two macro indicators were engineered:

1. **Total Income:** Ingests combined financial streams to represent actual macro repayment capacity:
$$Total_Income = ApplicantIncome + CoapplicantIncome.$$
2. **Income to Loan Ratio:** Calculates structural financial leverage to serve as an anchor feature against continuous variance:
$$Income_Loan_Ratio = \frac{Total_Income}{LoanAmount \times 1000}.$$

5. Model Building & Comparative Evaluation

Following strict data split validation rules (80% training matrix and 20% isolated test grid), three distinct machine learning algorithms were built, tuned, and cross-evaluated:

Model Architecture	Testing Accuracy	Precision Score	Recall Score	F1-Score
Logistic Regression	78.86%	0.7596	0.9875	0.8581
Decision Tree Classifier	68.29%	0.7468	0.7375	0.7421
Random Forest Classifier	78.04%	0.7731	0.9375	0.8474

Hyperparameter Optimization Insights: To maximize model efficiency, a parallelized `GridSearchCV` loop was deployed over the Random Forest architecture. Testing configurations across multi-tier tree depths and split parameters yielded an elite, stable cross-validation accuracy of **81.88%**, locking in optimal estimators.

6. Ensemble Architecture & Model Interpretability

To safeguard production cycles from variance anomalies, a consensus-driven meta-model—the **Voting Classifier**—was constructed. This combined the linear stability of Logistic Regression with the deep non-linear boundaries of the optimized Random Forest Classifier, maximizing system robustness.

Global feature importance tracking conclusively proved that `Credit_History` remains the ultimate gateway metric, followed directly by our engineered asset leverage feature, `Income_Loan_Ratio`.

7. Production Ingestion & Serialization Ready Engine

To prevent data leakage during real-time inference, data scaling parameters (`StandardScaler`) and the optimized hyperparameter weights were bound into an immutable `sklearn.pipeline.Pipeline` object.

The complete pipeline execution matrix has been compiled and serialized to disk as a binary artifact: `optimized_loan_predictor_pipeline.pkl` via `joblib`. Additionally, a native Python simulation engine function wrapper was developed to ingest live applicant JSON/dictionary parameters, perform automated feature synthesis on-the-fly, and return sub-second credit verdicts with absolute statistical confidence, making it instantly portable for full web API or Streamlit implementations.